

Cohort grading report

9 agents · 2 families · alphabetical order throughout

Canonical platform / measured truth	4 / 9
Non-canonical / flagged	5 / 9
Median improvement (all)	+0.0 %
Median improvement (canonical only)	+50.1 %
Range across cohort	+0.0 % ... +64.8 %

Families in this cohort

module-1	n = 4	module-2	n = 5
----------	-------	----------	-------


Contents

- Page 2 — Outcome bars: per-agent self-reported % improvement
- Page 3 — Attribution: variant contributions per agent
- Page 4 — Methodology rubric heatmap (agents × rubric items)
- Page 5 — Canonical re-evaluation by family (if canonical data available)


Caveats

Self-reported % uses each agent's chosen metric, unit, and split — not comparable across agents.
Canonical % (page 5) re-runs every model on the same fixed eval set — this IS comparable.
Grey bars on page 2 = agent scored on non-canonical platform or fabricated proxy.

Self-reported outcome — improvement on the lateral metric (alphabetical)



Attribution — variant contributions per agent (alphabetical)



Methodology rubric — green ✓ pass · red ✗ fail · grey — null / not addressed

m1-agent-01	✗	✓	✓	✗
m1-agent-02	✗	✗	✗	✗
m1-agent-04	✗	✓	✓	✗
m1-agent-05	✗	—	✗	—
m2-agent-01	✗	✗	✗	—
m2-agent-02	✗	✗	✗	✗
m2-agent-03	✗	✗	✗	✗
m2-agent-04	✗	✗	✗	—
m2-agent-05	✗	✗	✗	✓
	breakdown-present	methodology-consistent	attribution-coherent	self-regression-flagged