

Cohort grading report

10 agents · 2 families · alphabetical order throughout

Canonical platform / measured truth	9 / 10
Non-canonical / flagged	1 / 10
Median improvement (all)	+46.1 %
Median improvement (canonical only)	+47.0 %
Range across cohort	+0.0 % ... +64.8 %

Families in this cohort

module-1	n = 5	module-2	n = 5
----------	-------	----------	-------


Contents

- Page 2 — Outcome bars: per-agent self-reported % improvement
- Page 3 — Attribution: variant contributions per agent
- Page 4 — Methodology rubric heatmap (agents × rubric items)
- Page 5 — Canonical re-evaluation by family (if canonical data available)


Caveats

- Self-reported % uses each agent's chosen metric, unit, and split — not comparable across agents.
- Canonical % (page 5) re-runs every model on the same fixed eval set — this IS comparable.
- Grey bars on page 2 = agent scored on non-canonical platform or fabricated proxy.

Self-reported outcome — improvement on the lateral metric (alphabetical)



Attribution — variant contributions per agent (alphabetical)



Canonical evaluation by family — every model re-run on fixed Ford eval set

